

# Introduction to Hbase

Gkavresis Giorgos - 1470



# Agenda

- What is Hbase
- Installation
- About RDBMS
- Overview of Hbase
- Why Hbase instead of RDBMS
- Architecture of Hbase
- Hbase interface
- Summarise



# What is Hbase

Hbase is an open source, distributed sorted map modeled after Google's BigTable



# Open Source

- Apache 2.0 License
- Committers and contributors from diverse organizations like Facebook, Trend Micro etc.



# Installation

Download link

<http://www.apache.org/dyn/closer.cgi/hbase/>

Before starting it, you might want to edit

`conf/hbase-site.xml` and set the directory you want

HBase to write to, `hbase.rootdir`

Can be standalone or pseudo distributed and  
distributed

Start Hbase via `$ ./bin/start-hbase.sh`



# About Relational Database Management Systems

- Have a lot of Limitations
- Both read / write throughout not possible(transactional databases)
- Specialized Hardware is quite expensive



# Background

- Google releases paper on Bigtable – 2006
- First usable Hbase – 2007
- Hbase becomes Apache top-level project – 2010
- Hbase 0.26.5 released.





# Overview of Hbase

- Hbase is a part of Hadoop
- Apache Hadoop is an open-source system to reliably store and process data across many commodity computers
- Hbase and Hadoop are written in Java
- Hadoop provides:
  - Fault tolerance
  - Scalability





# Hadoop advantages

- Data parallel or compute-parallel. For example:
  - Extensive machine learning on <100 GB of image data
  - Simple SQL queries on >100 TB of clickstreaming data



# Hadoop's components

- MapReduce(Process)
  - Fault-tolerant distributed processing
- HDFS(store)
  - Self-healing
  - High-bandwidth
  - Clustered storage



# Difference Between Hadoop/HDFS and Hbase

HDFS is a distributed file system that is well suited for the storage of large files. HBase, on the other hand, is built on top of HDFS and provides fast record lookups (and updates) for large tables.

HDFS has based on GFS file system.



# Hbase is

- Distributed – uses HDFS for storage
- Column – Oriented
- Multi-Dimensional(Versions)
- Storage System



# Hbase is NOT

- A sql Database – No Joins, no query engine, no datatypes, no (damn) sql
- No Schema
- No DBA needed



# Storage Model

- Column – oriented database (column families)
- Table consists of Rows, each which has a primary key(row key)
- Each Row may have any number of columns
- Table schema only defines Column families(column family can have any number of columns)
- Each cell value has a timestamp



# Static Columns

int	varchar	int	varchar	int
int	varchar	int	varchar	int
int	varchar	int	varchar	int





# Something different


- Row1  $\rightarrow$  ColA = Value
- ColB = Value
- ColC = Value
- Row2  $\rightarrow$  ColX = Value
- ColY = Value
- ColZ = Value



# A Big Map

**Row Key + Column Key + timestamp  
=> value**

Row Key	Column Key	Timestamp	Value
1	Info:name	1273516197868	Sakis
1	Info:age	1273871824184	21
1	Info:sex	1273746281432	Male
2	Info:name	1273863723227	Themis
2	Info:name	1273973134238	Andreas



# One more example

Row Key	Data
cutting	Info: {'height': '9ft', 'state': 'CA'} Roles: {'ASF': 'Director', 'Hadoop': 'Founder'}
tlipcon	Info: {'height': '5ft7', 'state': 'CA'} Roles: {'Hadoop': 'Committer'@ts=2010 'Hadoop': 'PMC'@ts=2011 'Hive': 'Contributor'}



# Column Families

- Different sets of columns may have different priorities
- CFs stored separately on disk access one without wasting IO on the other.
- Configurable by column family
  - Compression(none,gzip,LZO)
  - Version retention policies
  - Cache priority




# Hbase vs RDBMS

	RDBMS	Hbase
Data layout	Row-oriented	Column family oriented
Query language	SQL	Get/put/scan/etc *
Security	Authentication/Authorization	Work in Progress
Max data size	TBs	Hundreds of PBs
Read / write throughput limits	1000s queries/second	Millions of queries per second



# Terms and Daemons

- Region
  - A subset of table's rows,
- RegionServer(slave)
  - Serves data for reads and writes
- Master
  - Responsible for coordinating the slaves
  - Assigns regions, detects failures of Region Servers
  - Control some admin function



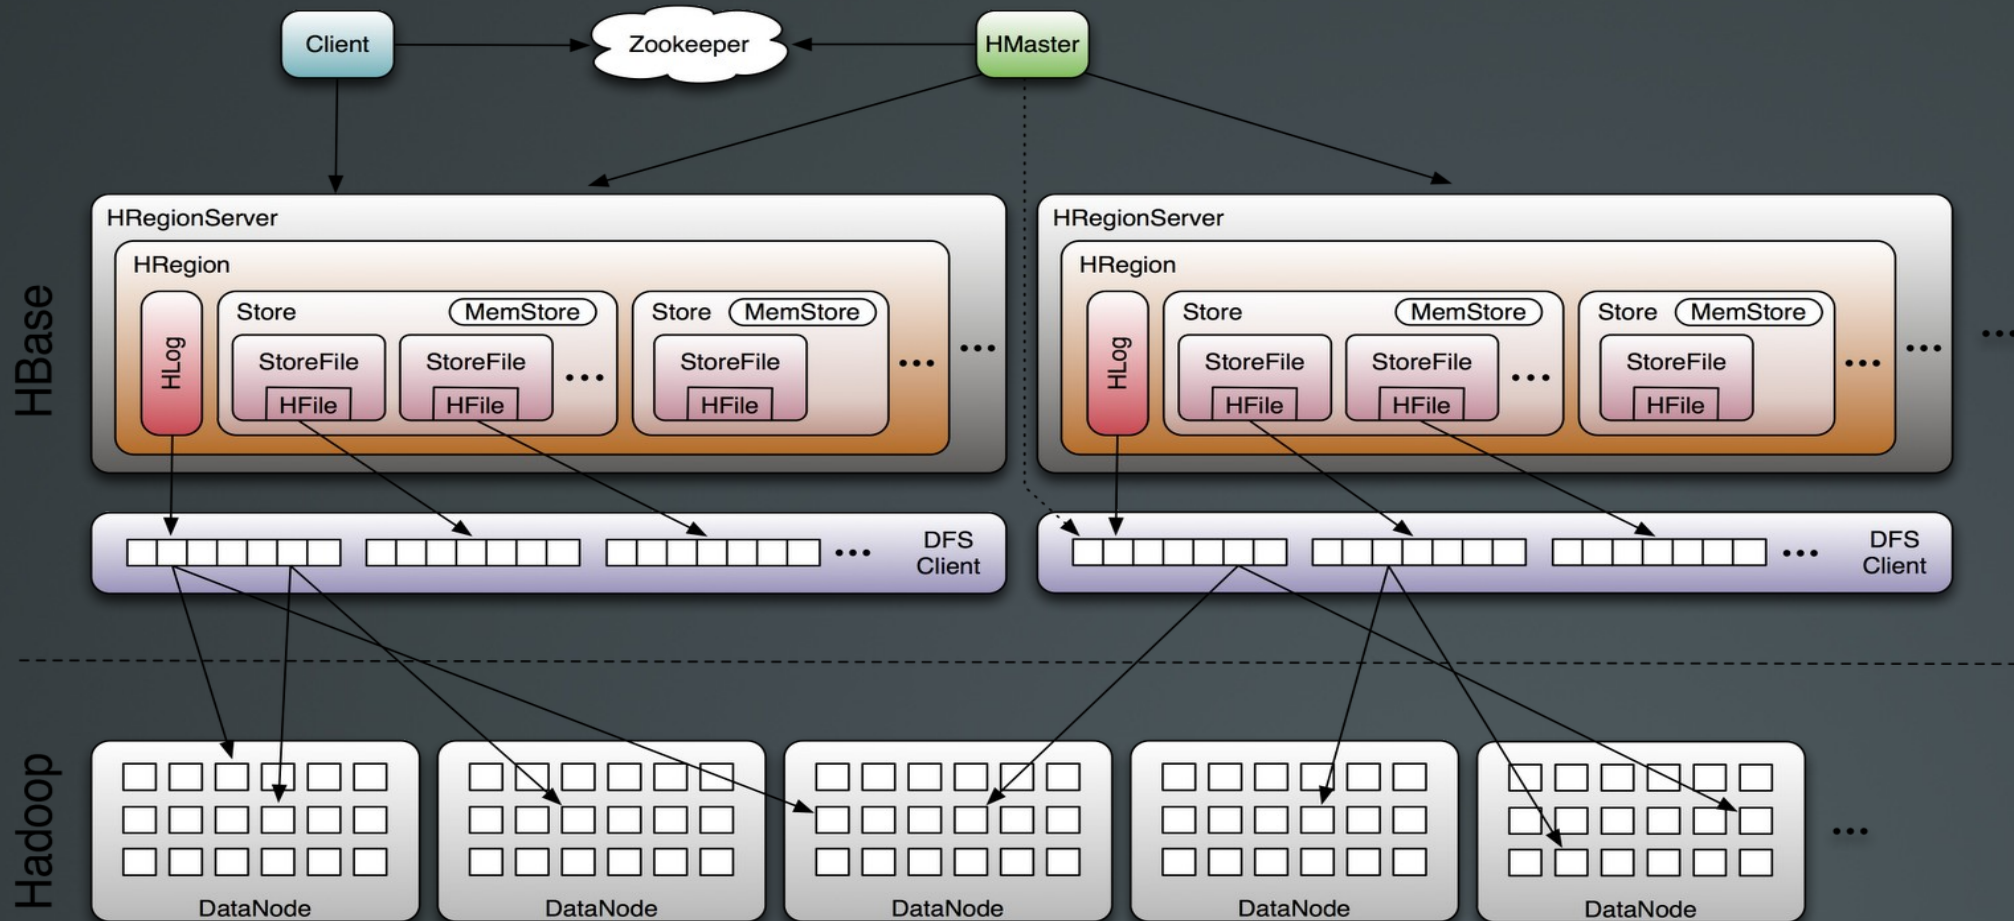
# Distributed coordination

- To manage master election and server availability we use Zookeeper
- Set up a cluster, provides distributed coordination primitives
- An excellent tool for building cluster management systems





# Hbase Architecture



# Distributed coordination

- To manage master election and server availability we use Zookeeper
- Set up a cluster, provides distributed coordination primitives
- An excellent tool for building cluster management systems



# Hbase Interface

- Java
- Thrift(Ruby,Php,Python,Perl,C++,...)
- Hbase Shell



# Hbase API

- `get(row)`
- `put(row, Map<column, value>)`
- `scan(key range, filter)`
- `increment(row, columns)`
- Check and Put, delete etc.



# Hbase shell

- `hbase(main):003:0> create 'test', 'cf'`
- 0 row(s) in 1.2200 seconds
- `hbase(main):004:0> put 'test', 'row1', 'cf:a', 'value1'`
- 0 row(s) in 0.0560 seconds
- `hbase(main):005:0> put 'test', 'row2', 'cf:b', 'value2'`
- 0 row(s) in 0.0370 seconds
- `hbase(main):006:0> put 'test', 'row3', 'cf:c', 'value3'`
- 0 row(s) in 0.0450 seconds



# Hbase shell cont.

- hbase(main):007:0> scan 'test'
- ROW       COLUMN+CELL
- row1      column=cf:a, timestamp=1288380727188, value=value1
- row2      column=cf:b, timestamp=1288380738440, value=value2
- row3      column=cf:c, timestamp=1288380747365, value=value3
- 3 row(s) in 0.0590 seconds





# Hbase in java

```
HBaseConfiguration conf = new HBaseConfiguration();  
conf.addResource(new Path("/opt/hbase-0.19.3/conf/hbase-site.xml"));  
  
HTable table = new HTable(conf, "test_table");  
BatchUpdate batchUpdate = new BatchUpdate("test_row1");  
  
batchUpdate.put("columnfamily:column1", Bytes.toBytes("some value")  
);  
  
batchUpdate.delete("column1");  
table.commit(batchUpdate);
```





# Get Data

Read one column value from a row

```
Cell cell = table.get("test_row1", "columnfamily1:column1");
```

To read one row with given columns, use `HTable#getRow()` method.

```
RowResult singleRow = table.getRow(Bytes.toBytes("test_row1")  
);
```



# A "tough" facebook application

- Realtime counters of URLs shared, links "liked", impressions generated
- 20 billion events/day (200K events/sec)
- ~30 sec latency from click to count
- Heavy use of incrementColumnValue API
- Tried MySQL, Cassandra, settled on Hbase



# Use Hbase if

- You need random write, random read or both (but not neither)
- You need to do many thousands of operations per sec on multiple TB of data
- Your access patterns are simple



Thank you \../

